

DSM Partitioning

A Pipeline for Reordering Dependency Structure Matrices

Simon Kraft Joshua Payne El Sall

CPSC 482/682 — Data Structures II
University of Northern British Columbia

April 7, 2026

Outline

- 1 Introduction
- 2 Methodology
- 3 Results
- 4 Conclusions

The challenge:

- Modern software & engineering projects involve **hundreds to thousands of interdependent tasks** across diverse teams
- Traditional tools (Gantt, CPM) handle sequential & parallel flows well...
- ...but **fail for interdependent tasks** where outputs feed back as inputs
- These feedback dependencies are common in complex product development

Goal

Find a **task ordering** that minimises conflicts in information flow to enable faster, more predictable project execution.

Key observation

When task j supplies input to task i but is scheduled *after* it, assumptions must be made. If wrong, task i must be **re-executed**.

Design Structure Matrix (DSM)

Definition: A binary matrix $A \in \{0, 1\}^{n \times n}$
where

$$a_{ij} = 1 \iff \text{task } i \text{ depends on task } j$$

- Rows and columns labeled by the same task list
- Diagonal entries carry no inter-task meaning
- Equivalent to directed graph
 $G(A) = (V, E)$

Three dependency types:

- **Feed-Forward** ($j < i$): below diagonal — easy
- **Feed-Back** ($j > i$): above diagonal — problematic
- **Coupled**: mutual dependency — hardest

The Partitioning Problem

Objective: Find $P \in \{0, 1\}^{n \times n}$ such that $A' = P^T A P$ minimises:

Objective 1 — Feed-Back Marks

$$|\text{FBM}| = |\{a'_{ij} \neq 0 \mid j > i\}|$$

Minimise the **count** of above-diagonal nonzeros.

Objective 2 — Feedback Distance

$$\text{TFBD} = \sum_{a'_{ij} \in \text{FBM}} (j - i)$$

Minimise the **distance** of feed-back marks from the diagonal.


Best case

$G(A)$ has no directed cycles $\Rightarrow$ fully lower triangular,
 $|\text{FBM}| = 0$, $\text{TFBD} = 0$

General case

Achieve **Block Lower Triangular Form (BLTF)**: group cycle-participants into diagonal blocks and minimise feedback distance.

Our 6-Stage Pipeline



Data Structure: CSC

- Column pointer array `col_ptr[0...n]`
- Row index array `row_ind[0...nnz-1]`
- Storage: $O(n + \text{nnz})$
- No value array needed (binary DSM)

Why CSC?

Column j lists out-neighbours of vertex j , giving $O(\deg^+(j))$ access — exactly the pattern required for DFS-based SCC algorithms.

SCC Decomposition

Both algorithms implemented as **iterative DFS** over CSC, using an explicit stack (no recursion).

Tarjan's Algorithm

- **Single DFS pass**
- Each vertex v maintains $low(v)$: smallest reachable pre-order time
- $low(v) = v.pre$ at finish $\Rightarrow v$ is SCC root; pop stack
- $O(V + E)$ time, $O(V)$ space

Kosaraju-Sharir Algorithm

- **Two DFS passes**
- Pass 1: DFS on G , record finish order
- Pass 2: DFS on $rev(G)$ in reverse finish order; each DFS tree = one SCC
- $O(V + E)$ time, $O(V + E)$ space

After SCC decomposition, the **condensation DAG** is topologically sorted with **Kahn's algorithm** to determine the block ordering.

Tearing & Banding

Tearing (intra-block FBM reduction)

- Approximates the **minimum feedback arc set** within each non-trivial SCC
- Uses **Algorithm GR** (Eades et al., 1993): greedily removes sinks, then sources, then the vertex with maximum $\delta = d^+ - d^-$
- Backward edges in the returned ordering are torn
- $O(k^2 + m)$ per SCC (linear scan variant)

Banding (TFBD reduction)

- **Reverse Cuthill-McKee (RCM)**: concentrates nonzeros near diagonal
- Pseudoperipheral start via George-Liu BFS
- **Adjacent-swap refinement**: accepts swaps that decrease TFBD without increasing FBM; up to 20 passes
- Best of torn + RCM orderings is kept

Benchmark data

10 sparse matrices from the project repository, $n = 11$ to 27,770, $\text{nnz} = 4$ to 352,807. Compiled

Stage-wise Quality Improvement

Average reductions across 10 matrices:

Stage	FBM ↓	TFBD ↓
Baseline	—	—
TopoOnly	60.16%	86.56%
+ Tearing	76.47%	<i>slight</i> ↑
+ Banding	76.50%	85.95%

- Topo ordering provides the **largest single gain**
- Tearing drives further FBM reduction within SCCs
- Banding polishes TFBD without hurting FBM

Selected matrix results

Matrix	FBM ₀ → FBM _f	TFBD _f
Tina_AskCal	12 → 4	7
California	6063 → 192	669
wb-cs-stanford	20131 → 8537	1.78M
HEP-th-new	104435 → 8281	10.1M

Runtime Comparison: Tarjan vs. Kosaraju-Sharir

Key findings

- Both algorithms produce **identical SCC decompositions**
- Tarjan: mean $\approx 560 \mu\text{s}$
- Kosaraju-Sharir: mean $\approx 992 \mu\text{s}$
- Tarjan is $\approx 1.77\times$ faster on average
- Scaling compatible with expected $O(V + E)$ near-linear behaviour

Pipeline stage runtimes

Stage	Relative cost
Tarjan SCC	low
Kosaraju SCC	$1.77\times$ Tarjan
Condensation	low
Topo sort	low
Tearing	moderate
Banding	highest
Permutation	low

Why is Tarjan faster?

Single DFS pass vs. two passes + transpose graph construction in Kosaraju-Sharir.

Banding is the heaviest stage but delivers the final quality gains.

Qualitative Results — Matrix Spy Plots

California ($n = 9,664$)

Near-perfect partitioning: almost no residual FBM, no large SCC blocks visible.

wb-cs-stanford ($n = 9,914$)

One large persistent SCC block retains significant FBM far from diagonal — explains high residual TFBD.

Tina_AskCal ($n = 1,023$)

Clean block structure with well-separated components after reordering.

GD99_c ($n = 62$)

Illustrates the SCC structure clearly: many trivial singleton and two-node SCCs, plus two larger coupled blocks.

(See appendix figures for full spy plot comparisons before and after partitioning.)

Conclusions & Future Work

What we achieved

- Implemented full 6-stage DSM reordering pipeline in C++ with CSC
- Both Tarjan and Kosaraju-Sharir SCC algorithms — identical decompositions
- **76.50% average FBM reduction, 85.95% average TFBD reduction** across 10 benchmark matrices
- Tarjan $1.77\times$ faster than Kosaraju-Sharir
- Runtime scales near-linearly with $V + E$

Future directions

- **Better tearing:** replace greedy Algorithm GR with an exact or improved heuristic for minimum feedback arc set
- **Parallelism:** per-SCC tearing & banding are independent — parallelise for large matrices
- **Weighted DSMs:** extend to numerical dependency strengths; optimise weighted feedback cost

Acknowledgements & Team Contributions

AI Disclosure

Claude (Anthropic) and ChatGPT (OpenAI) were used to aid understanding of concepts and sourcing literature.

Team Member Contributions

Member	Contribution
Simon Kraft	[placeholder]
Joshua Payne	[placeholder]
El Sall	[placeholder]

Thank you — questions welcome!